

DATA
SCIENCE
CLUB

Intro to Pytorch

Data Science Club

Agenda

- Introduction to PyTorch
- Introduction to Tensors / PyTorch basics
- How PyTorch supports a ML Workflow
- Extensions of NN's
- Extra Tools in PyTorch

Introduction to PyTorch

- PyTorch is an Open-source deep learning framework
- Developed by Facebook's AI Research lab
- Strengths of PyTorch
 - Pythonic syntax
 - Ability to build both simple and complex models
 - Large and active community
 - Supports Automatic Differentiation for Neural networks

PyTorch v.s other Frameworks

- Comparison with TensorFlow:
 - TensorFlow uses static computation graphs (older versions)
 - PyTorch uses dynamic computation graphs (easier for debugging)
 - Both support large-scale machine learning deployments
- Why PyTorch?
 - More flexibility (especially for R&D and experimentation)
 - Natural integration with Python libraries (e.g., NumPy)
 - Simple, readable syntax

PyTorch Fundamentals: Tensors

What are Tensors?

- Multi-dimensional arrays (similar to NumPy arrays)
- Generalization of vectors and matrices
- Used to represent input data and parameters in deep learning

PyTorch Tensor vs NumPy Array:

- PyTorch tensors support GPU acceleration
- PyTorch tensors come with autograd support
- Can easily convert between NumPy arrays and PyTorch tensors

```
import torch  
x = torch.tensor([1, 2, 3])  
print(x)
```

```
tensor([1, 2, 3])
```

Tensor Operations

- **Element-wise Operations:**
 - Operations like addition, subtraction, multiplication, and division are performed element-wise between tensors of the same shape.
- **Matrix Multiplication:**
 - PyTorch supports matrix multiplication using either the `@` operator or `torch.matmul()` function.
- **Reshaping Tensors:**
 - You can change the shape of tensors using `.view()` or `.reshape()`, which is especially useful when preparing data for neural networks.
- **Slicing and Indexing:**
 - Similar to NumPy, tensors in PyTorch can be sliced and indexed.

Tensor Operations: Sample code

```
import torch

x = torch.tensor([1.0, 2.0, 3.0])
y = torch.rand(1, 3) # Random tensor with values sampled from [0, 1)

print("X: ", x)
print("Y: ", y)

z = torch.zeros(3, 3) # 3x3 matrix of zeros
print("Z: ", z)

result = torch.add(x, y) # x + y
print("Result of add: ", result)
result = torch.mul(x, y) # x * y
print("Result of mul: ", result)

x = torch.randn(2, 3)
print("ran X: ", x)
x_resaped = x.view(3, 2) # Reshaping the tensor
print("reshaped tensor", x_resaped)

x[:, 1] # All rows, second column
print("sliced tensor", x[:, 1])

X: tensor([1., 2., 3.])
Y: tensor([[0.6740, 0.8895, 0.6841]])
Z: tensor([[0., 0., 0.],
          [0., 0., 0.],
          [0., 0., 0.]])
Result of add: tensor([[1.6740, 2.8895, 3.6841]])
Result of mul: tensor([[0.6740, 1.7790, 2.0523]])
ran X: tensor([[ -2.1624, -0.5272,  0.9708],
               [-0.3355, -1.9378, -0.9678]])
reshaped tensor tensor([[ -2.1624, -0.5272],
                        [ 0.9708, -0.3355],
                        [-1.9378, -0.9678]])
sliced tensor tensor([-0.5272, -1.9378])
```

Data Loading with PyTorch

Efficient Data Loading:

- Neural networks require training on large datasets. PyTorch provides the `torch.utils.data.DataLoader` class to simplify data loading and manage mini-batches during training.

Dataset and DataLoader:

- **`torch.utils.data.Dataset`**: Represents your data and how to access it. You can create a custom dataset by subclassing this class and overriding `__len__` and `__getitem__` methods.
- **`torch.utils.data.DataLoader`**: Handles batching, shuffling, and loading data in parallel using workers (multiprocessing).

Batch Size and Shuffling:

- **Batch Size**: Controls the number of samples per batch passed to the model.
- **Shuffling**: Randomizes the order of data samples, which helps to prevent the model from learning data order instead of patterns.


```
from torch.utils.data import DataLoader

train_loader = DataLoader(dataset=train_dataset, batch_size=64, shuffle=True)
for batch in train_loader:
    inputs, labels = batch |
```



PyTorch and Neural Networks

First steps: Logistic Regression

Remember how logistic models work: They have a set of parameters (b_0, b_1) and output a probability corresponding to a certain data input (x).



Neural networks take this idea and extend it.




Next step: A Neural Network

We can instead have multiple logistic regressions in parallel that each have different values for b .

Now each one of these outputs can be used as the data input for a next logistic regression.


During the training process the model learns to diversify each set of b values in a way that gives more useful information.




1 layer with 3 neurons

Example of logistic regression v.s NN's

Logistic Regression



Tuned Neural Network (hidden=10, activation=logistic)



Introduction to Neural Networks

Introduction to Neural Networks

- **What are Neural Networks?**
 - Neural networks are models inspired by the human brain that consist of interconnected layers of neurons.
 - Each neuron in a neural network performs a simple mathematical operation, and the layers are stacked to learn complex patterns.
 - In PyTorch, neural networks are typically represented as a sequence of layers (fully connected layers, convolutional layers, etc.) followed by activation functions (ReLU, sigmoid, etc.).
- **PyTorch's `torch.nn` Module:**
 - PyTorch provides the `torch.nn` module for building neural networks. This module includes layers like `Linear` for fully connected layers and functions like `ReLU` for non-linear activations.
 - Each neural network is defined as a class that subclasses `nn.Module` and implements a `forward`

Building Neural Networks in PyTorch

Defining a Neural Network:

- In PyTorch, neural networks are typically defined as Python classes that inherit from `torch.nn.Module`. This class acts as a container for layers and implements the forward pass logic.
- The core structure involves two steps:
 1. **Initializing Layers:** In the `__init__()` method, define the layers (e.g., fully connected layers, convolutions) that will make up your network.
 2. **Forward Method:** In the `forward()` method, define how data moves through the layers in your model. This is where the layers you've defined are applied to the input data, layer by layer.

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(784, 128) # Fully connected layer (input: 784, output: 128)
        self.fc2 = nn.Linear(128, 10)  # Fully connected layer (input: 128, output: 10)

    def forward(self, x):
        x = F.relu(self.fc1(x)) # Apply ReLU activation after the first layer
        x = self.fc2(x)         # Output layer, no activation (usually combined with a loss function)
        return x
```

Training a Neural Network

Training Loop in PyTorch:

- PyTorch uses a simple, clear training loop that follows the steps:
 - Forward Pass:** Send the input through the model to get the predicted output.
 - Compute Loss:** Use a loss function to calculate the difference between the predicted output and the ground truth (target labels).
 - Zero Gradients:** Reset the gradients before the backward pass by calling `optimizer.zero_grad()`. Otherwise, gradients accumulate.
 - Backward Pass:** Call `loss.backward()` to compute the gradients of the loss with respect to each parameter.
 - Update Weights:** Use `optimizer.step()` to update the parameters using the computed gradients.

```
for epoch in range(num_epochs):  
    outputs = model(inputs) # Forward pass  
    loss = criterion(outputs, labels) # Compute loss  
  
    optimizer.zero_grad() # Clear gradients  
    loss.backward() # Backward pass (compute gradients)  
    optimizer.step() # Update weights
```

Loss Function in PyTorch

What are Loss Functions?

- Loss functions measure how well a model's predictions match the true labels. They provide feedback to adjust model weights during training.
- PyTorch offers a variety of built-in loss functions in the `torch.nn` module.

Common Loss Functions:

- **Classification Tasks:**
 - Use `CrossEntropyLoss` for multi-class classification problems, which combines `LogSoftmax` and `NLLLoss`.
- **Regression Tasks:**
 - Use `MSELoss` (Mean Squared Error) for regression tasks where the target is a continuous variable.
- **Other Losses:**
 - **L1 Loss:** Measures the absolute differences between predicted and target values.
 - **Binary Cross-Entropy (BCE):** Commonly used for binary classification.
- **Choosing the Right Loss:**
 - Depending on your task, selecting the appropriate loss function is crucial to guide the learning process effectively.

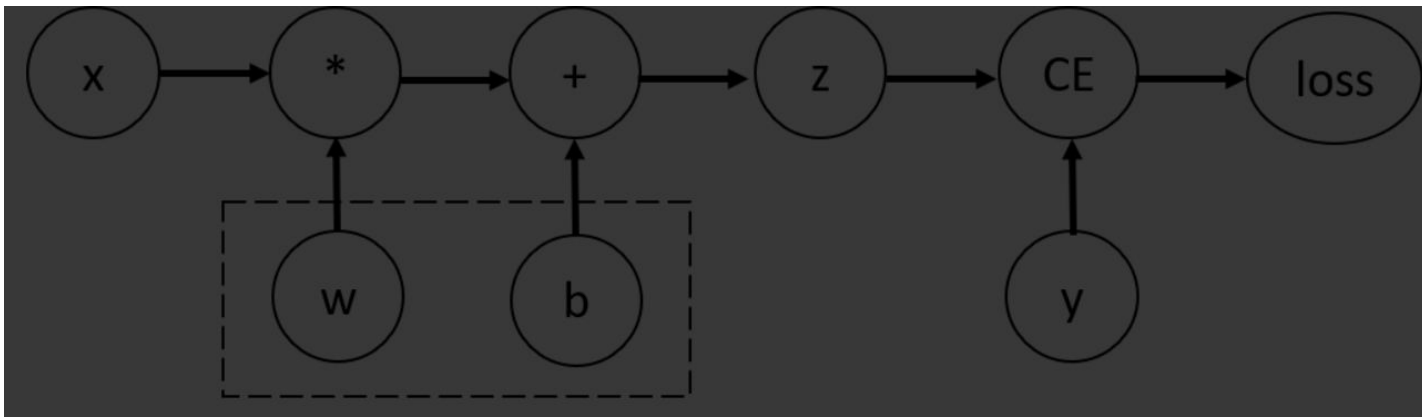
Backpropagation in PyTorch

Backpropagation Overview:

- Core to how neural networks learn
- Calculates gradients of the loss function with respect to model parameters

PyTorch's `backward()` Function:

- Computes the gradient of the loss function
- Automatically populates the `.grad` attribute for each tensor



Autograd Automatic Differentiation

What is Autograd?

- PyTorch's `autograd` module automatically computes gradients of tensor operations. These gradients are crucial for training neural networks because they indicate how to adjust model parameters to minimize the loss function.
- PyTorch tracks operations on tensors with `requires_grad=True` and constructs a computational graph. When the `.backward()` function is called, the framework traverses the graph to calculate gradients.

Why is this important?

- This automation is key to the backpropagation algorithm, which is fundamental in deep learning to adjust weights based on the loss.

```
import torch

x = torch.ones(2, 2, requires_grad=True) # Track gradients
y = x + 2 # Perform operations
z = y * y * 3 # Further operations
z.backward() # Compute gradients
print(x.grad) # Gradients stored in x.grad

tensor([[18., 18.],
        [18., 18.]])
```

Optimizers in PyTorch

Role of Optimizers:

- After calculating gradients via backpropagation, an optimizer adjusts the model's parameters to minimize the loss function.
- Optimizers use the computed gradients to update the model's weights. PyTorch provides several optimization algorithms in the `torch.optim` module.

Popular Optimizers:

- **SGD (Stochastic Gradient Descent):**
 - One of the simplest and most commonly used optimizers. SGD updates model parameters with learning rate and momentum:
- **Adam (Adaptive Moment Estimation)**
 - A popular choice because it adapts the learning rate for each parameter based on gradient moments. It often works well with minimal tuning.
- **RMSprop:**
 - A variant of SGD that scales the learning rate based on recent gradients, commonly used in RNNs
- **Learning Rate:**
 - The learning rate controls how big of a step the optimizer takes during each parameter update. It's a crucial hyperparameter that needs tuning.

```
optimizer = torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
```

Training on GPU with PyTorch

Why Use GPUs?

- GPUs are designed for fast matrix operations, making them significantly faster than CPUs for training deep neural networks. PyTorch supports GPU acceleration using CUDA.

Using CUDA in PyTorch:

- **GPU-accelerated Training:** All tensor operations performed on GPU are much faster than on CPU, especially when dealing with large datasets and models.

```
#Checking for GPU availability
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

#Moving model and data to GPU
model.to(device)
inputs, labels = inputs.to(device), labels.to(device)
```

Extensions of NN's: Convolutional Neural Networks

CNNs are particularly well-suited for tasks involving:

- Image recognition
- Natural language processing
- Time series analysis

Imagine a sliding window.

- **Input:** A signal (like an image or audio).
- **Window:** A smaller, movable pattern (often called a kernel or filter).

Why is it useful?

- **Feature extraction:** CNNs use convolutions to learn and extract meaningful features from input data.
- **Pattern recognition:** By sliding the filter across the input, CNNs can detect patterns and structures.
- **Noise reduction:** Convolution can be used to smooth out noise in signals.

Transfer Learning in PyTorch

- **Why Save Models?**
 - Saving a trained model allows you to reuse it without retraining, share it with others, or deploy it in production.
- **Saving Model Weights:**
 - PyTorch provides a simple way to save model parameters using the `torch.save()` function. This saves the model's state dictionary, which contains the weights and biases.
- **Loading Model Weights:**
 - To reload the saved weights, first create an instance of the model, then load the state dictionary.
- **Saving and Loading Entire Models:**
 - It's also possible to save the entire model, including its architecture and weights. This can be useful for quick deployment.

Transfer Learning in PyTorch: sample code

```
#Saving model weights
torch.save(model.state_dict(), 'model_weights.pth')

#Loading model weights
model = Net()
model.load_state_dict(torch.load('model_weights.pth'))

#Saving/Loading entire models
torch.save(model, 'model.pth')
model = torch.load('model.pth')
```

Link to Colab exercise

- Link: <https://tinyurl.com/mry2zk9e>
 - Be sure to make a copy of the notebook before you start!
- Slides: <https://tinyurl.com/dscptworkshopslides>

Free Perplexity for 1 year!!!

(scan QR code and create account
with cmu account)

